Lecture 4

Goal: Multi-programming

run N programs on M processors
(N >> M)

suppose m = 1

multiplexing - time sharing.

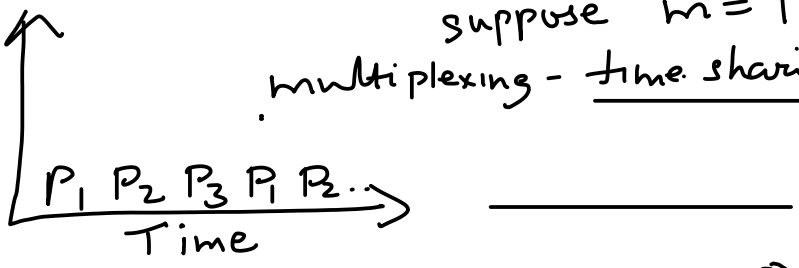

Recall:

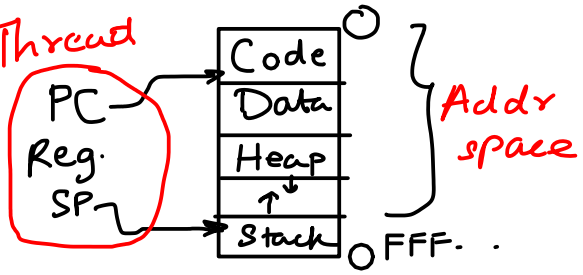

* What is the benefit of multi-programming?

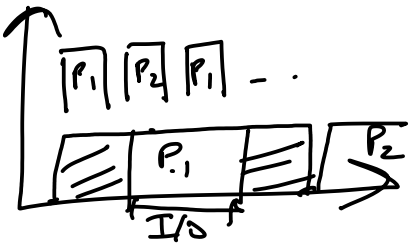

→ Improve utilization of resources

I/O - read a file block

Speed of CPU >> I/O op.

→ Interactive system

- handle interrupts

(*) The OS performs the switch from 1 prog to another prog.
- {Save the running prog's state & load the state of other prog.}

Important Considerations!

(1) Efficiency

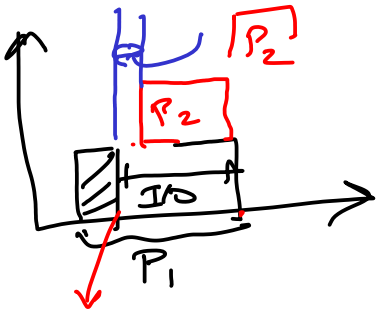

Switching time ≃ I/O

switching time → 0

(2) Safety

Switching from 1 prog. to another should not corrupt the state of any prog/ OS.

Which component performs the switch?
- S/W or (H/W)

Protection

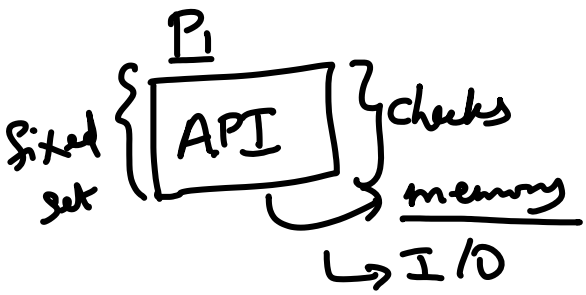

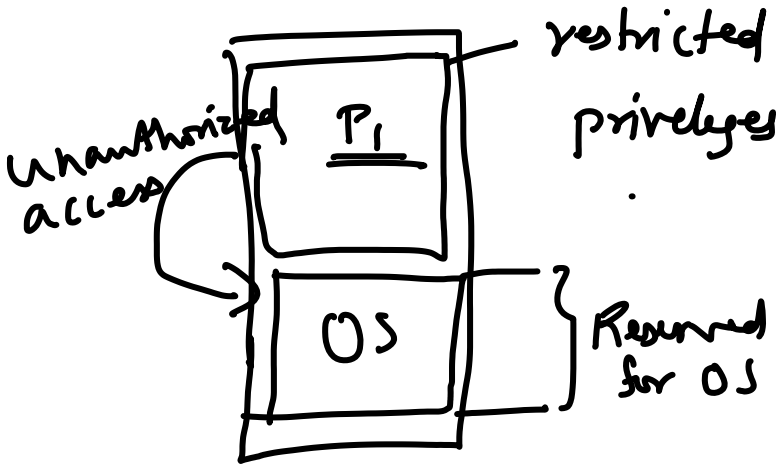

Process Abstraction (*) 2.2 | 2.3

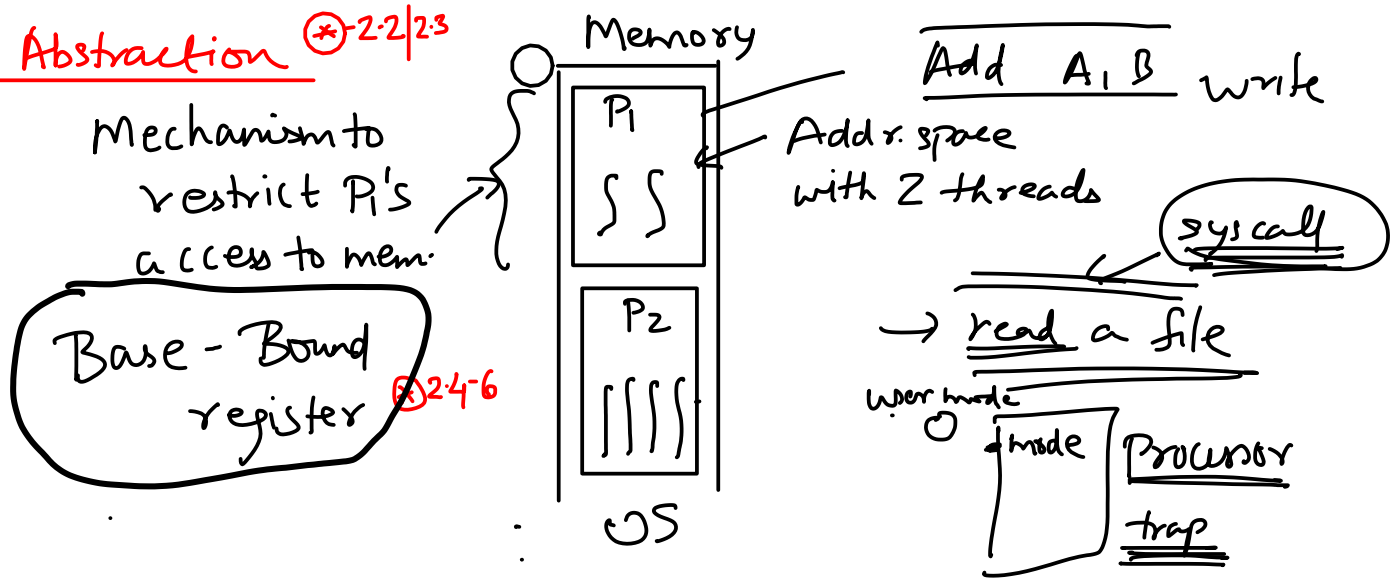

Process – **Restricted privelges** OS – unrestricted privelges

Issues to be addressed

(I) Process wants to do something restricted ?
- eg. I/O read a file

Soln: (*) 2.7-9

different privelge levels + mechanism to change from 1 privelge level to another } sys. call, interrupts, trap, exceptions

dual-mode of operation

e.g. (*) 2.10-13

fork()

Inst. Trap

1) saves state of running program
   - done by h/w
2) change privilge level from user - kernel
3) run the correct trap handler

PC - user code → syscall

Kernel mode { PC - prog. counter
OS-kernel.

How do you find the sys-call handler? } User prog

uPC:

Kernel boots

- Kernel mode
- Initializes sys-call tables

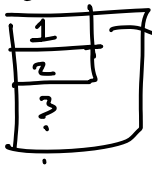

→ defn - open (- Check for access to the file
- read the block into kernel mem
- RTU

RTU - uPC → PC
1 → 0
User mode

∴ user mode : opn

restricted opn - trap → user to kernel
- trap handler

Recap

Multi-programming

efficiency Safety

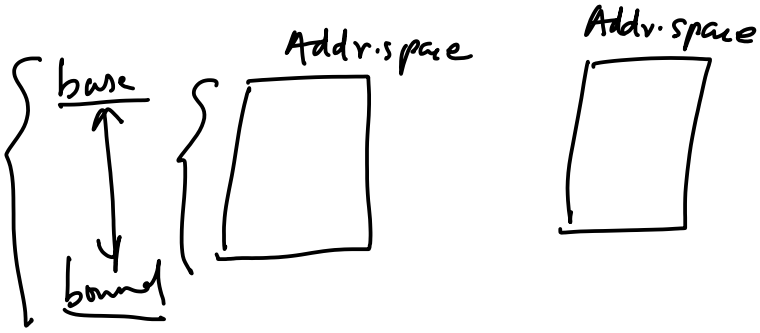

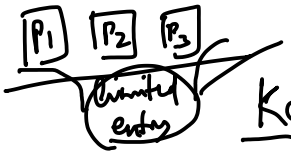

Kernel mode

Protection: – OS provides the process abstraction

↳ Privelged inst. – non-priveleged inst. → (Addr.space)

↳ Memory protection → user mode

↳ Timer interrupts

Kernel

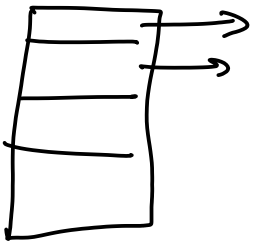

divide by zero user mode

↓ exception

→ kernel mode

→ saves the prg state

→ handler –

illusion

exceptions Privelged

gdb – breakpoints

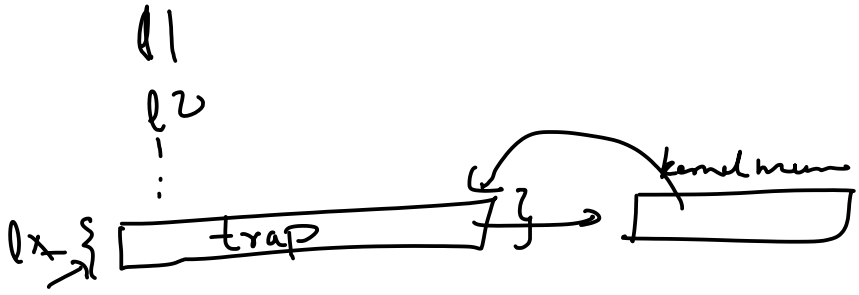

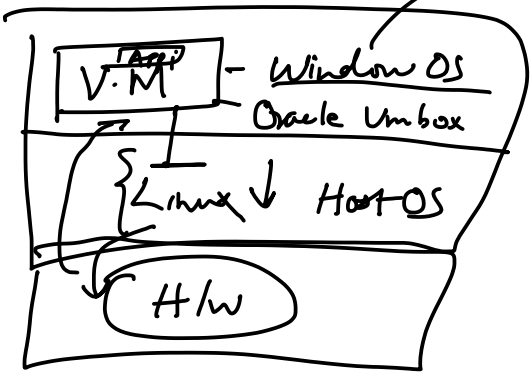

(11) Context switching - how does OS stop 1 process & load another process?

(*) 2.14 - 19

Limited entry pt.

syscalls - read, open,

#syscall

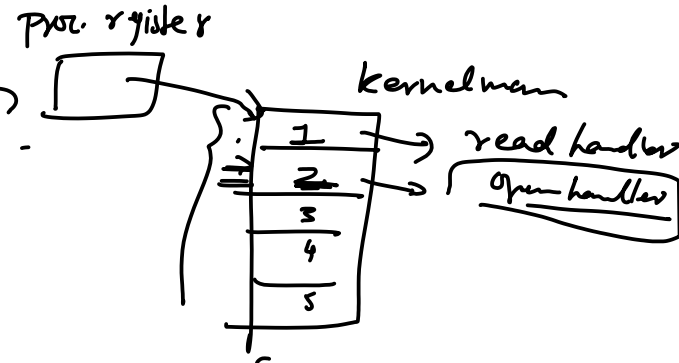

user stack

- kernel stack

```
main() {
  func(2,3)
}
func(a, b) {
```

Stack - func(a,b)

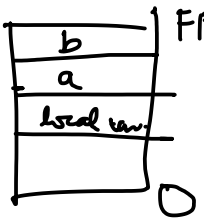

sys.call - file_open() -

}

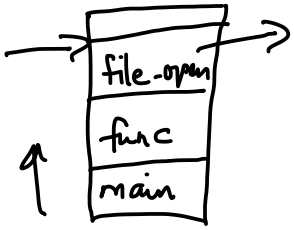

Kernel should NOT user user-level stack

Security { (1) SP may be invalid
(2) stack can be manipulated by some thread

wait for I/O

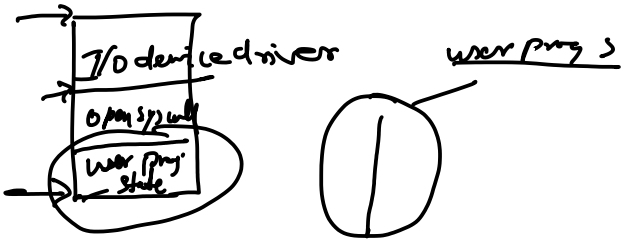

Kernel interrupt stack per process

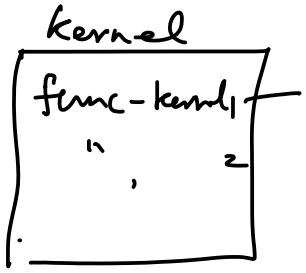

user stub

file_open( ) {

push #Sys call

trap

}

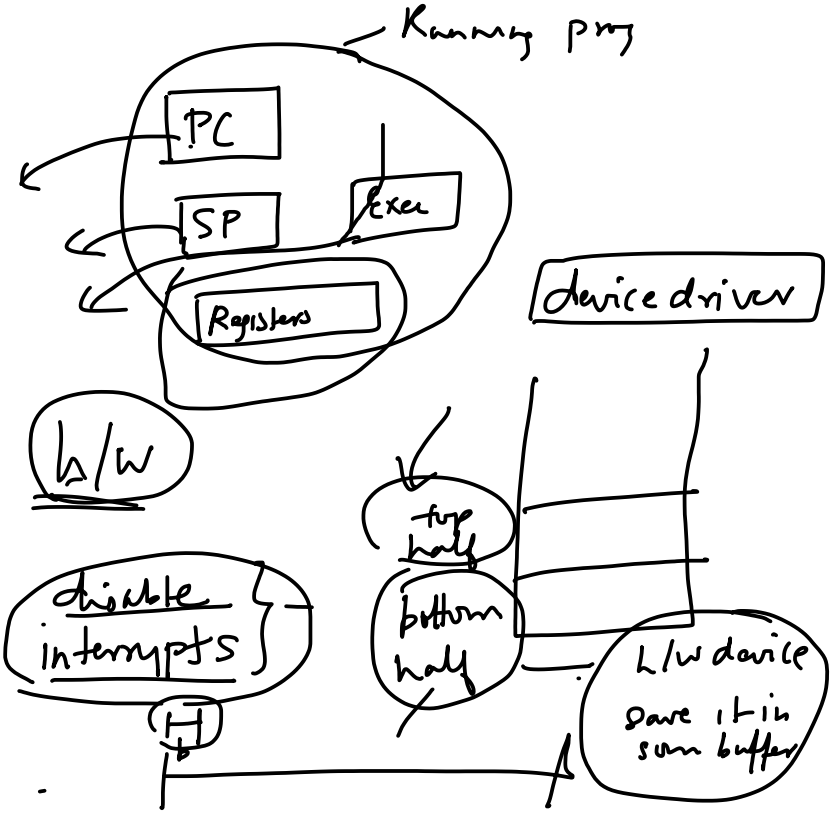

Interrupt

Servicing Interrupt

another interrupt arrives

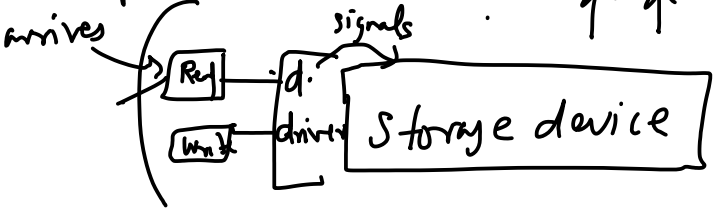

Multiprogramming → improve utilization

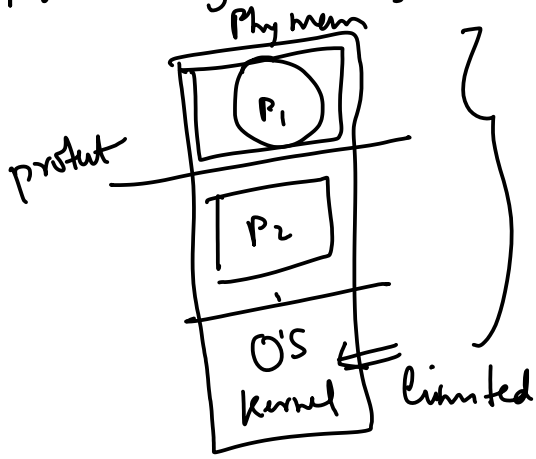

— protection

↳ priveleged inst.

↳ based bound mem protection

↳ Timer Interrupt

2002 - Apple's OS

reboot the OS

X { yield( ) program — voluntarily relinquishes its CPU

Timer Interrupt → OS has to decide which proc. to run

$P_1 \xrightarrow{I/O}$ I/O not complete

State of a process

fork() → creating a new process

1 processor
- 200 processes
- 1 process running
- 199

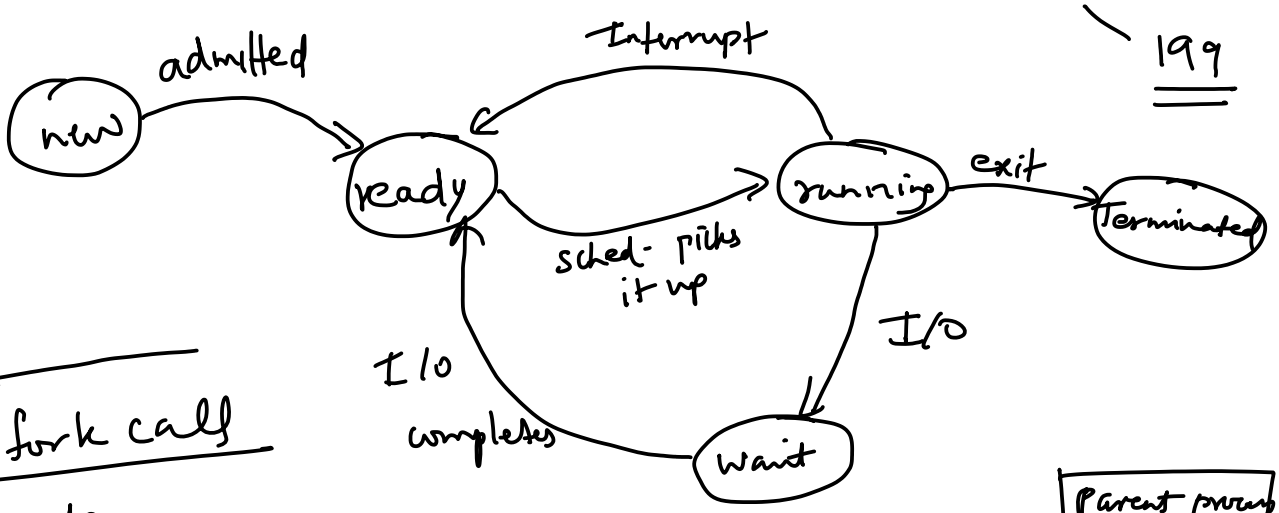

On fork call
- Create PCB_1 in the kernel

- create/initialize the addr space
- inherits the exec context of the parent
- Add the PCB_1 to the ready list

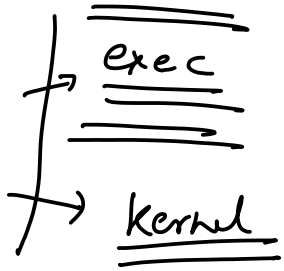

Shell — ls

PCB — Pid, user, privileges, open fd, status — OS Kernel
↳ Kernel memory

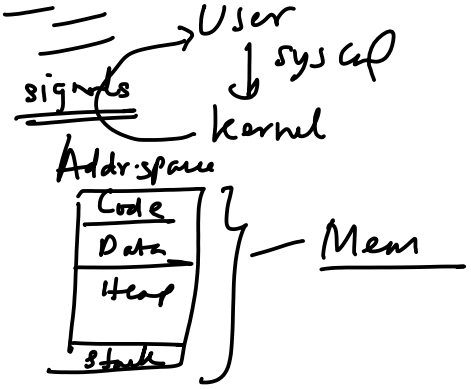

Zombie process

id=fork();

if (id==0)

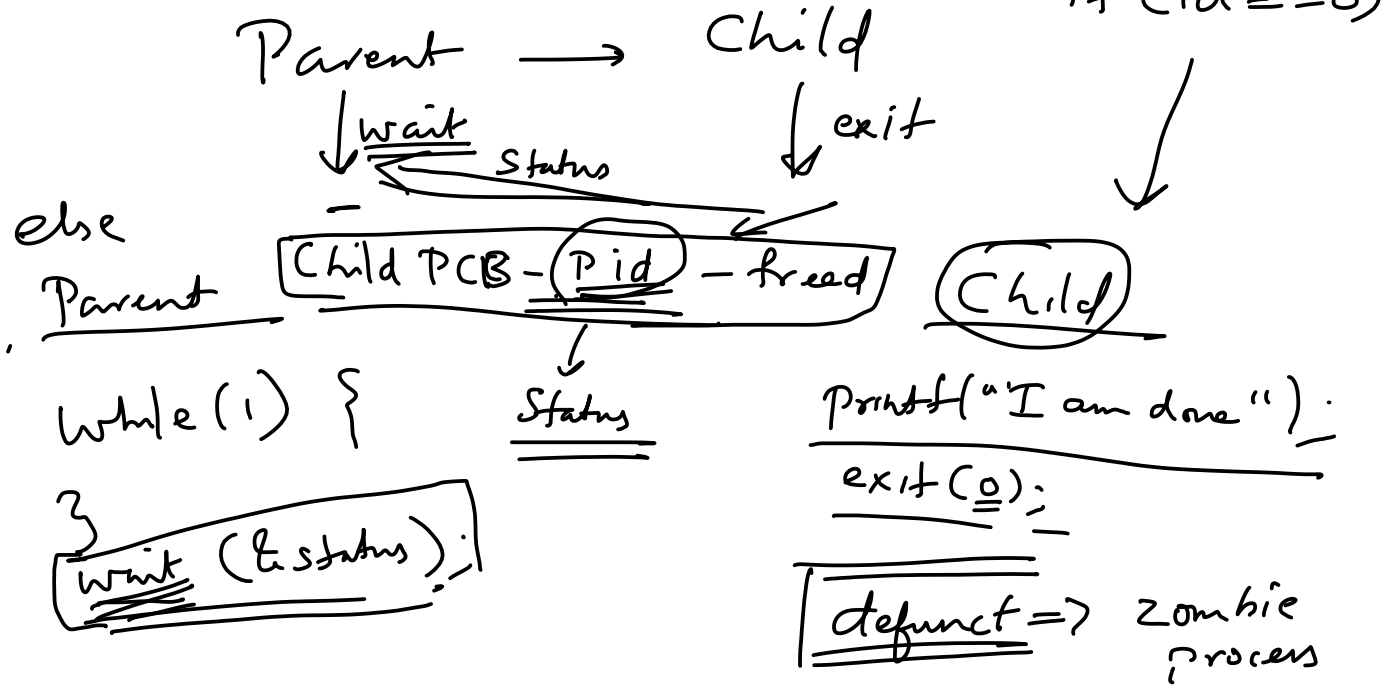

fork bomb

```
while(1)
  fork()
```

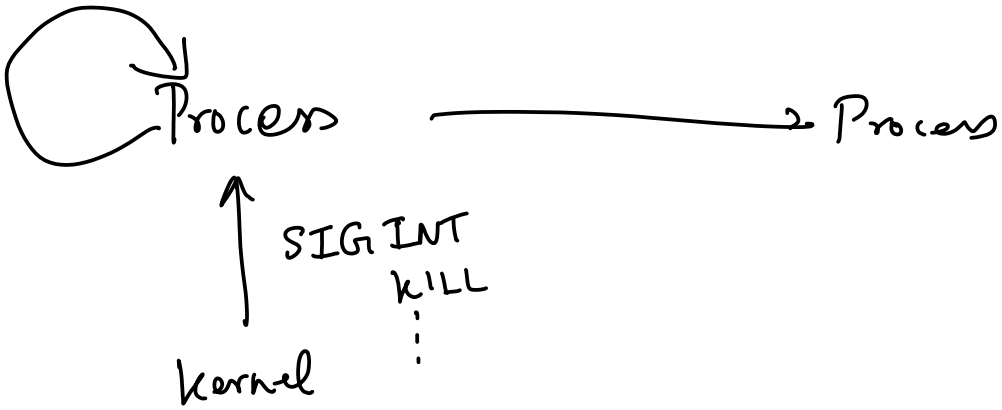

Database

Concurrency - threads (T_1, T_2, T_3 :)

Database

Own thread library

/API

thread Create

:

Signals

OS - Thread library

API - thread Create

:

Kernel threads

1s 2>

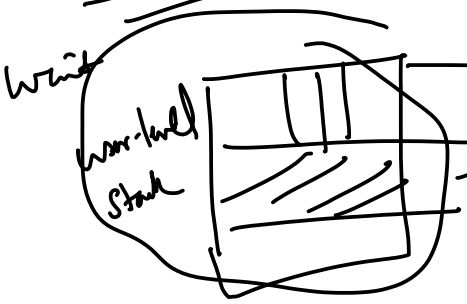

T_1 - Syscall read file block

T_2

PCB

TCB_1 . TCB_2

P

user-level thrads

P

T_1 T_2

Kernel Kernel

Kernel interrupt stack