


# Concurrency

Process: Addr. space with 1 or more threads




Why do we need threads?



→ Overlap I/O operations with Compute  $op^n$  —

→ Multiple processors — 1 thread cannot exploit the parallelism

Single-threaded



T1	T2
f1 = (a, b, c)	f2 (d, e, f, g)

-d
C

TCB: Thread Control Block

PCB<sub>1</sub> — PCB<sub>2</sub> — /proc / pid / tid<sub>1</sub>  
 ↳ TCB<sub>1</sub> ↳ TCB<sub>2</sub>

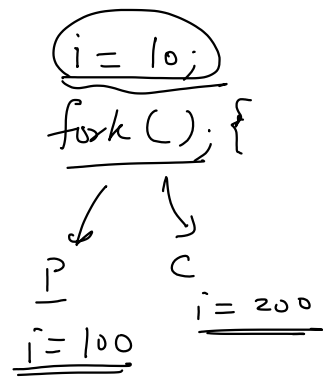

T1 shares some data with T2 — Sharing with threads is easy

P1 share some data with P2 —

Inter-process comm<sup>n</sup>

Pipes

∴ give rise to sync. issues.



Issues that arise because of sharing state

{ non-reproducible  
non-deterministic }

↳ hard to debug programs with threads

Q.  $\frac{\text{arg} = 200}{\text{arg} = 20000000}$  — o/p — 400  
Non-deterministic! } Context-switching

Fundamental building block

atomic operations —  $\text{bal} = \text{bal} + 1$  } atomic op<sup>n</sup>  
without any interruption { load bal } CS {  $\text{lock.acquire()}$  }  $\text{lock.acquire()}$   
store }  $\text{bal} = \text{bal} + 1$  } wait  
 $\text{lock.release()}$


2 threads cannot hold the same lock at the same time.

{ Critical section }  
mutual exclusion

time: → real — exec. time  
user — time spent in user mode  
sys — " in kernel mode

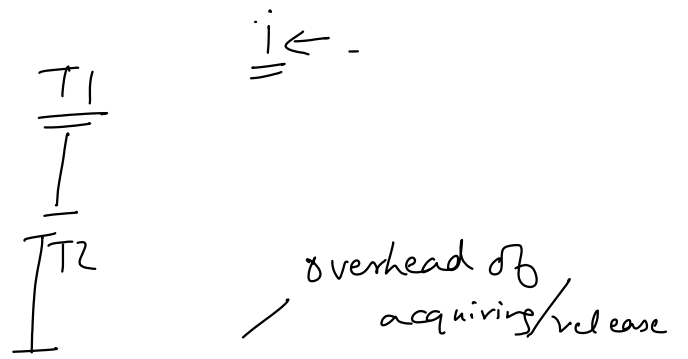
[ user < real  
sys < real ]

parallelism  
P1 P2  
T1 T2  
[ { T1 } { T2 } ]  
real = max(T1, T2)  
[ user = T1.user + T2.user ]



fine-grained locks —

⇓  
 { more room for  
 parallelism }



Coarse-grained lock

⇓  
degree of parallelism ↓